

I tested 4 AI models on 137 messy business documents.


**The best one still got 11% of them
wrong somewhere.**

1980s faxed invoices. Receipts from Jakarta and
Kuala Lumpur. US tax forms. Indian bank statements.
Real SEC contracts.


Strict scoring: a field is right only if
the value is right. Swipe →

Field accuracy hides the problem.

One wrong field means a person re-checks the whole document. So I count documents with EVERY field right.



Whole document right, by type



* synthetic. n = 10 to 32 per type, so gaps of a few points are noise.

GPT-5.6 Terra 'fixes' names.

e Employee's first name and initial	Last name	Suff.
Rachael	Bates	
282 Monique Trail Apt. 902 East Thomas, CA 95553		

Claude Opus 5.5

Rachael

Claude Sonnet 5

Rachael

Qwen3-VL 8B (laptop)

Rachael

GPT-5.6 Terra

Rachel

Also: Kelleyland → Kellyland, Stevenmouth →
Steventown, 'and Sons' → 'and Songs'.
On a tax form, each one is a wrong name.

The free laptop model:

**beat GPT-5.6 Terra on W-2s (21/32 vs 7/32),
then failed Indian bank statements (2/10).**

Printed: **02-01-2024**

Model read: **1 Feb 2024**

Correct: **2 Jan 2024**

Every amount and every balance was right.
Only the dates flipped, the American way.
That is the whole problem for Indian accounting.

**'Check your own work'
changed nothing.**

119 of 137

answers came back identical when GPT-5.6 Terra was shown the document again and asked to correct its own extraction.

It fixed 1 document and broke 1.
A second pass is not an evaluation.

The model bill is noise.

GPT-5.6 Terra	~\$5.60 per 1,000 docs (measured)
Claude Sonnet 5	~\$4.50 (estimated)
Claude Opus 5.5	~\$9.80 (estimated)
Qwen3-VL 8B, laptop	\$0

Under 1 cent a document for every model.

The errors are the cost: a wrong SSN,
a flipped date, a 'corrected' name.

Method and limits

- Same prompt for every model, full-resolution images
- Claude ran with no tools; Qwen ran locally via Ollama
- Scorer tested on 31 cases, including ones that must fail
- 10 to 32 documents per type: small gaps are noise
- W-2s and bank statements are synthetic; the W-2s were generated this week, so no model has seen them
- Invoice answer keys hand-transcribed, then checked by a second person against the images
- Public datasets may be in training data

Prompts, answer keys and every raw output: link in the comments.

Next: fine-tuning the free 8B model to fix these failures. Results in two weeks, win or lose.